

# Scapegoat Trees

## AAPP 2017 Project



Analysis, C implementation and testing, C++ benchmarks

Team members:

Crippa Mattia - 10397252

Vetere Alessandro - 10425802

# Introduction

- Scapegoat trees are **self balancing binary search trees**
- No extra data in the tree nodes
- Scapegoat scheme:
  - Tree:
    - Root node pointer
    - Total size
    - Maximum total size since last completely rebuilt
  - Node:
    - Key value
    - Left and right child nodes pointers
- Rebalancing operation:
  - On INSERT
  - On DELETE

# Fundamental properties

- $\alpha$ -weight-balanced:

$$\text{size}(\text{node.left}) \leq \alpha * \text{size}(\text{node}) \quad \wedge \quad \text{size}(\text{node.right}) \leq \alpha * \text{size}(\text{node})$$

- $\alpha$ -height-balanced,  $n = \text{size}(\text{root})$ :

$$h(T) \leq h_{\alpha}(n) = \lfloor \log_{1/\alpha} n \rfloor$$

- If  $T$  is an  $\alpha$ -weight-balanced binary search tree, then  $T$  is  $\alpha$ -height-balanced
- Scapegoat trees are **loosely**  $\alpha$ -height-balanced,  $n = \text{size}(\text{root})$ :


$$h(T) \leq h_{\alpha}(n) + 1$$

- Fixed  $\alpha$ , a node of depth greater than  $h_{\alpha}(n)$  is a **deep** node


# Meaning of $\alpha$

- Scapegoat node is an  $\alpha$ -weight-unbalanced node used to perform a **rebalance** operation
- Subtree rooted at the scapegoat node becomes an  $\frac{1}{2}$ -weight-balanced tree
- Value of  $\alpha$  should be in  $[0.5, 1)$ :
  - High  $\alpha \rightarrow$  Few rebalance operations, quick INSERT, but slow SEARCH and DELETE
  - Low  $\alpha \rightarrow$  Many rebalance operations, quick SEARCH and DELETE, but slow INSERT
- Choose  $\alpha$  depending on expected frequency of actions

# Maximum tree depth given $\alpha$



# What happens as $\alpha$ approaches 1



# SEARCH

- SEARCH proceeds as in a BST
- No restructuring is performed
- Due to loosely  $\alpha$ -height-balanced property, **worst case** time:

$$O(h_q(n)) = \mathbf{O(\log n)}$$

# INSERT

- INSERT a node as into a BST, update size and max\_size and record depth of new node
- If new node is **deep**, rebalance the tree:
  - Find scapegoat node
  - Rebalance the subtree rooted at the scapegoat
- Rebalancing operation will restore the  $\alpha$ -height-balanced property
- Worst case time  $O(\text{size}(\text{scapegoat}))$ , but **amortized** time:

**$O(\log n)$**




# How to choose scapegoat node


- Climb the tree from the new node back up to the root and select the first node that isn't  $\alpha$ -weight-balanced
- A better heuristic is choose the first node  $x_i$  satisfying the following condition, being  $i$  the minimum distance of  $x_i$  from the new node:

$$i > h_{\alpha}(\text{size}(x_i))$$


## INSERT Example (1/5)



## INSERT Example (2/5)




## INSERT Example (3/5)



## INSERT Example (4/5)



## INSERT Example (5/5)



$$h_a(n) = 5$$

# DELETE

- DELETE proceeds as in a BST
- Uses **max\_size** in the rebuilding condition:

$$\text{size} < \alpha * \text{max\_size}$$

- Worst case time  $O(n)$ , but **amortized** time:

$$\mathbf{O(\log n)}$$

## Pros

- SEARCH **worst case** time  **$O(\log n)$**
- INSERT and DELETE **amortized** time  **$O(\log n)$**
- No additional per node information

## Cons

- Two extra values per tree
- $\alpha$  should be tuned for expected frequency of operations



# Implementation in C

- Fast, robust and concise code
- Rich debug output at runtime if compiled using flag `-DDEBUG`
- Multiplatform, tested on Windows 10 64 bit and macOS Sierra
- Uses few external libraries:
  - `stdio.h` → `printf`
  - `math.h` → `log`, `floor`, `ceil`
  - `stdlib.h` → `malloc`, `free`
- Almost pedantic w.r.t. implementation described in the paper

# C API

```
typedef struct sg_node {  
    int key;  
    struct sg_node* left;  
    struct sg_node* right;  
} t_sg_node;
```

```
typedef struct sg_tree {  
    t_sg_node* root;  
    unsigned int size;  
    unsigned int max_size;  
    double alpha;  
    double log_one_over_alpha;  
    unsigned int h_alpha;  
} t_sg_tree;
```

```
t_sg_tree* sg_create_tree(double alpha);
```

```
void sg_delete_tree(t_sg_tree* tree);
```

```
t_sg_node* sg_search(t_sg_tree* tree, int key);
```

```
unsigned char sg_delete(t_sg_tree* tree, int key);
```

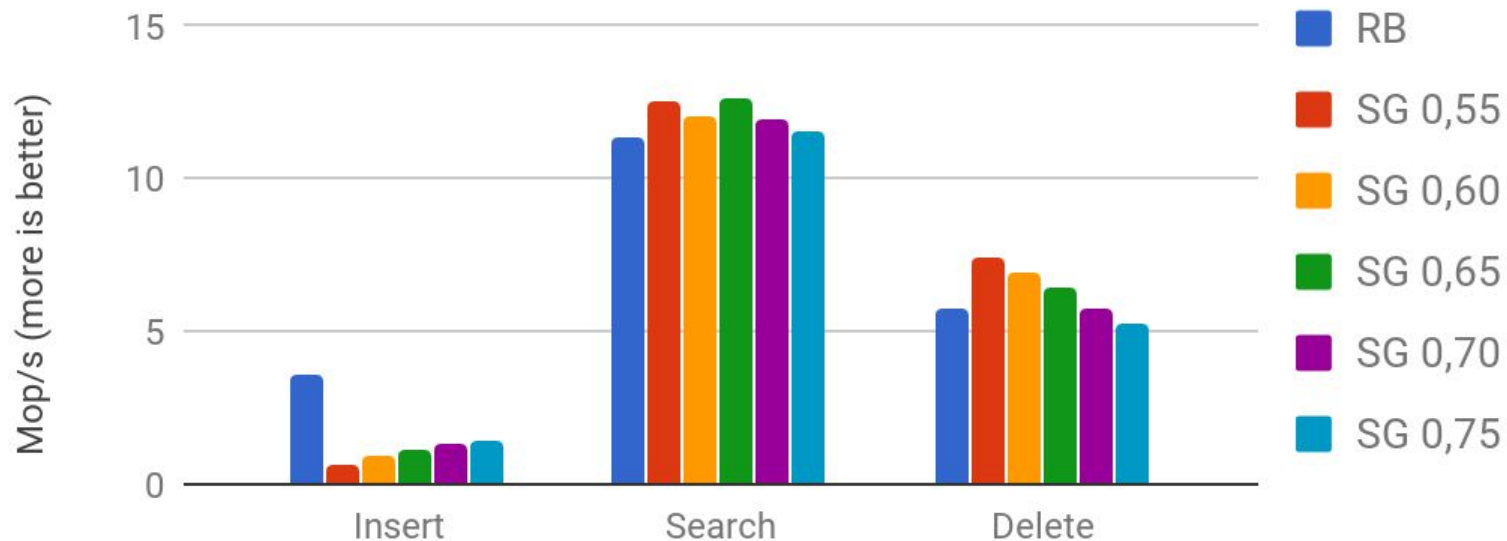
```
unsigned char sg_insert(t_sg_tree* tree, int key);
```

```
void sg_clear_tree(t_sg_tree* tree);
```

# Benchmark in C++ (1/4)


- Versus `std::set<int>` → red-black tree

SEQUENTIAL DATA 10 \* 3 Mop



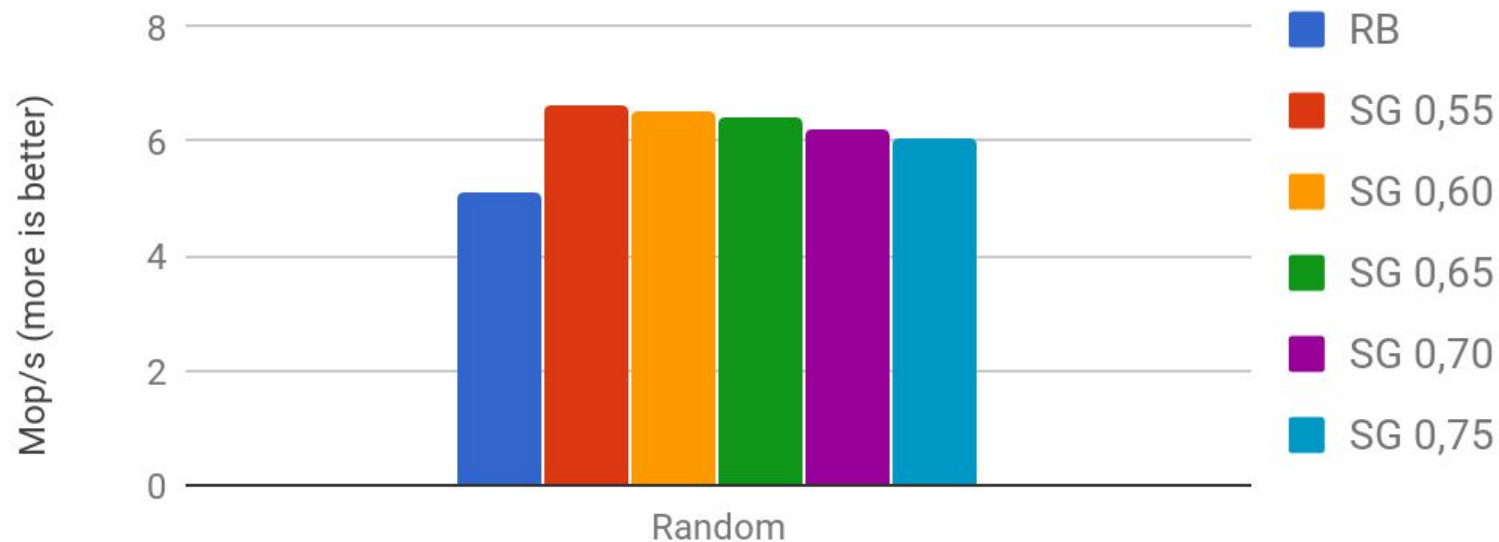
## Benchmark in C++ (2/4)

SHUFFLED DATA 10 \* 3 Mop




## Benchmark in C++ (3/4)

RANDOM 10k keys 10 Mop



## Benchmark in C++ (4/4)

MEMORY USAGE 10M keys



# Conclusions

- Sequential and shuffled INSERT faster on **red-black tree**, due to faster rebalancing operation
- Sequential SEARCH comparable, the two schemes yield balanced trees after sequential INSERTs
- Shuffled SEARCH much faster on **scapegoat tree**
- Sequential and shuffled DELETE comparable, a bit faster on **scapegoat tree** due to sporadic but complete rebalancing of the tree
- Random operations on random keys are faster on **scapegoat tree**, which demonstrates its strength as a general purpose ordered data structure
- **Scapegoat tree** is much more memory efficient than **red-black tree**